



KV Cache Pilot | Qwen2.5-1.5B

RTX 3090 · BF16 · eager · batch 1 · prompt 128 · generation 128 · greedy

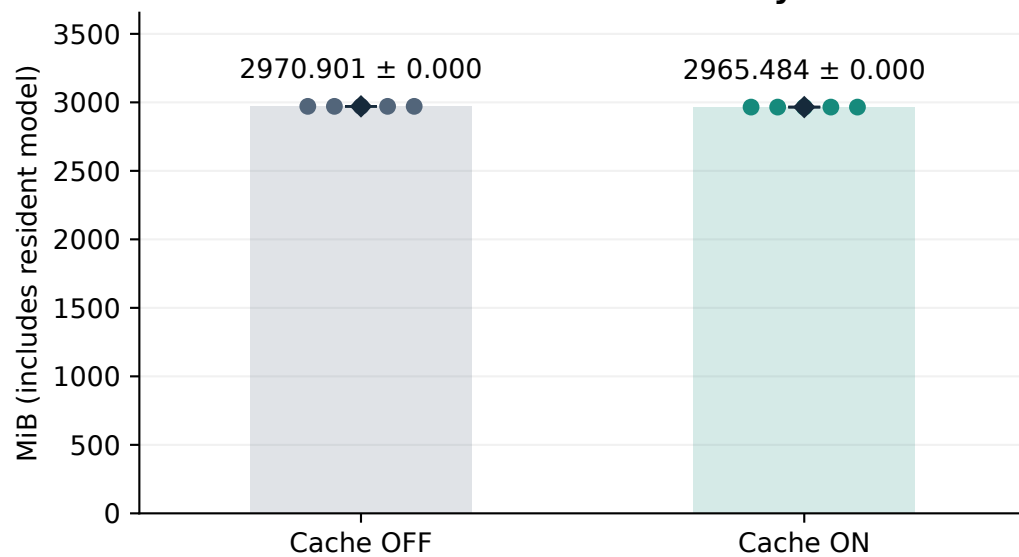
Total generation latency



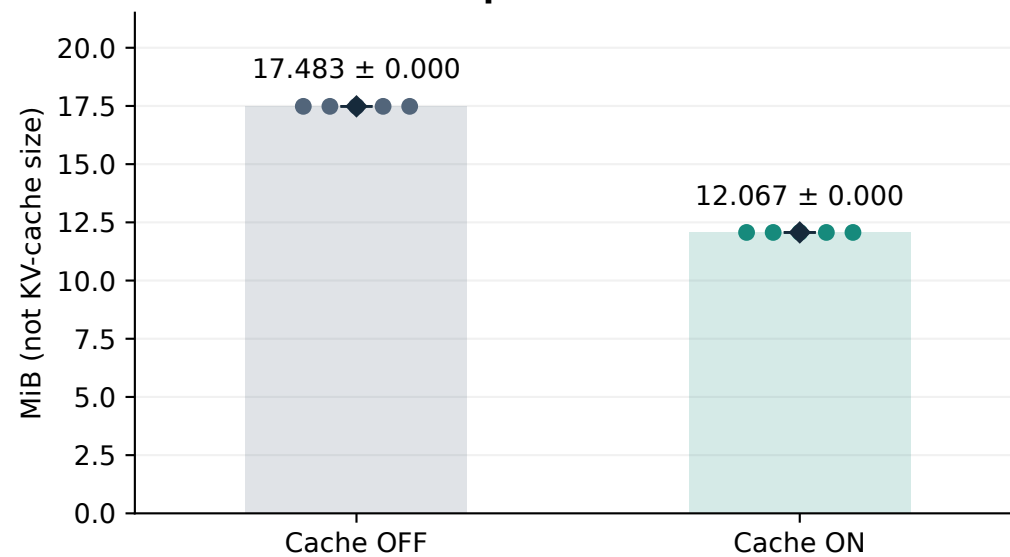
Generated-token throughput



Peak allocated memory



Allocated peak above baseline



Dots: 5 measured runs per mode. Diamond/error bars: mean ± sample SD (ddof=1). All axes start at zero.

Both modes: allocated baseline 2953.42 MiB; peak reserved 3184 MiB. Warm-up excluded; outputs identical.

Latency speedup (OFF mean / ON mean): 0.9922×. No speed gain observed in this pilot; descriptive results only.